

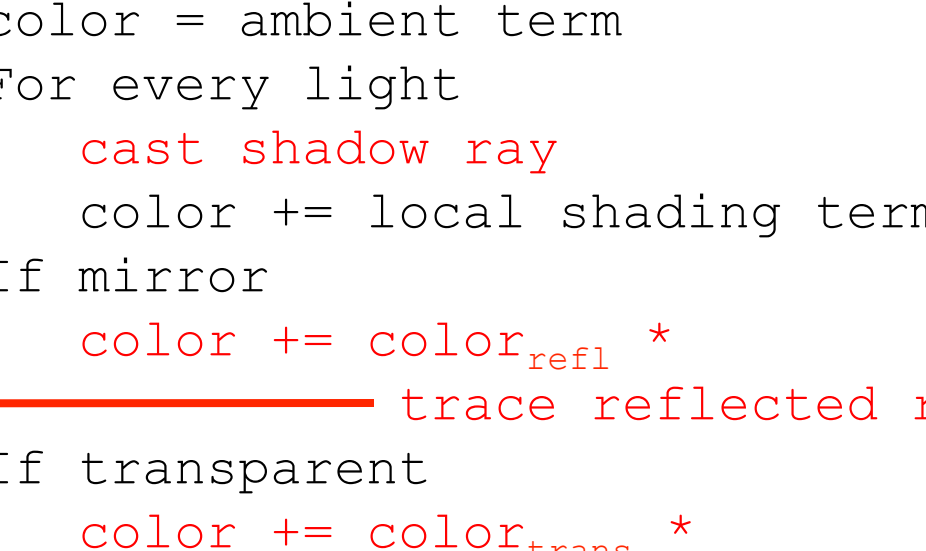
Acceleration Structures for Ray Casting

A 3D rendered scene of an ancient Egyptian temple interior. The scene features several columns with papyrus capitals and hieroglyphs on the walls. The architecture is made of light-colored stone. The lighting is warm, suggesting an indoor environment with light coming from the right. The columns are arranged in a row, with a doorway visible in the background.

ISTD 50.017 Graphics & Visualization
Sai-Kit Yeung, ISTD SUTD

Recap: Ray Tracing


trace ray



```
graph TD; A[Intersect all objects] --> B[color = ambient term]; B --> C[For every light]; C --> D[cast shadow ray]; D --> E[color += local shading term]; E --> F[If mirror]; F --> G[color += color_refl *]; G --> H[trace reflected ray]; H --> I[If transparent]; I --> J[color += color_trans *]; J --> K[trace transmitted ray];
```

Intersect all objects
color = ambient term
For every light
 cast shadow ray
 color += local shading term
If mirror
 color += color_{refl} *
 trace reflected ray
If transparent
 color += color_{trans} *
 trace transmitted ray


- Does it ever end?




Stopping criteria:

- Recursion depth
 - Stop after a number of bounces
- Ray contribution
 - Stop if reflected / transmitted contribution becomes too small


Recursion For Reflection: None



Recursion For Reflection: 1




Recursion For Reflection: 2



Ray tree


- Visualizing the ray tree for single image pixel



Ray tree

This gets pretty complicated
pretty fast!

- Visualizing the ray tree for single image pixel



Questions?

Ray Tracing Algorithm Analysis

- Lots of primitives
- Recursive
- Distributed Ray Tracing
 - Means using many rays for non-ideal/non-pointlike phenomena
 - Soft shadows
 - Anti-aliasing
 - Glossy reflection
 - Motion blur
 - Depth of field

cost \approx height * width *
num primitives *
intersection cost *
size of recursive ray tree *
num shadow rays *
num supersamples *
num glossy rays *
num temporal samples *
num aperture samples *
...

Can we reduce this?

Today

- Motivation
 - You need LOTS of rays to generate nice pictures
 - Intersecting every ray with every primitive becomes the bottleneck
- Bounding volumes
- Bounding Volume Hierarchies, Kd-trees

For every pixel

Construct a ray from the eye

For every object in the scene


Find intersection with the ray

Keep if closest

Shade


Accelerating Ray Casting

- Goal: Reduce the number of ray/primitive intersections




Conservative Bounding Volume

- First check for an intersection with a conservative bounding volume
- Early reject: If ray doesn't hit volume, it doesn't hit the triangles!






Conservative Bounding Volume

- What does “conservative” mean?
 - Volume must be big enough to contain all geometry within




Conservative Bounding Regions

- Desiderata
 - Tight → avoid false positives
 - Fast to intersect




Ray-Box Intersection

- Axis-aligned box
- Box: $(X_1, Y_1, Z_1) \rightarrow (X_2, Y_2, Z_2)$
- Ray: $P(t) = R_o + tR_d$




Naïve Ray-Box Intersection

- 6 plane equations: Compute all intersections
- Return closest intersection inside the box
 - Verify intersections are on the correct side of each plane: $Ax+By+Cz+D < 0$




Reducing Total Computation

- Pairs of planes have the same normal
- Normals have only one non-zero component
- Do computations one dimension at a time



Test if Parallel


- If $R_{dx} = 0$ (ray is parallel) AND $R_{ox} < X_1$ or $R_{ox} > X_2 \rightarrow$ no intersection



(The same for Y and Z, of course)


Find Intersections Per Dimension

- Basic idea
 - Determine an interval along the ray for each dimension
 - The intersect these 1D intervals (remember CSG!)
 - Done!




Find Intersections Per Dimension

- Basic idea
 - Determine an interval along the ray for each dimension
 - The intersect these 1D intervals (remember CSG!)
 - Done!



Find Intersections Per Dimension

- Basic idea
 - Determine an interval along the ray for each dimension
 - The intersect these 1D intervals (remember CSG!)
 - Done!




Interval
between X_1
and X_2

Interval
between Y_1
and Y_2

Find Intersections Per Dimension

- Basic idea
 - Determine an interval along the ray for each dimension
 - The intersect these 1D intervals (remember CSG!)
 - Done!




Interval
between X_1
and X_2

Interval
between Y_1
and Y_2


Intersection

Intersecting 1D Intervals




Intersecting 1D Intervals

Start=
max of mins




Intersecting 1D Intervals




Intersecting 1D Intervals

If $\text{Start} > \text{End}$, the intersection is empty!




Find Intersections Per Dimension

- Calculate intersection distance t_1 and t_2




Find Intersections Per Dimension

- Calculate intersection distance t_1 and t_2
 - $t_1 = (X_1 - R_{ox}) / R_{dx}$
 - $t_2 = (X_2 - R_{ox}) / R_{dx}$
 - $[t_1, t_2]$ is the X interval




Then Intersect Intervals

- Init t_{start} & t_{end} with X interval
- Update t_{start} & t_{end} for each subsequent dimension




Then Intersect Intervals

- Compute t_1 and t_2 for Y...




Then Intersect Intervals

- Update t_{start} & t_{end} for each subsequent dimension
 - If $t_1 > t_{\text{start}}$, $t_{\text{start}} = t_1$
 - If $t_2 < t_{\text{end}}$, $t_{\text{end}} = t_2$




Then Intersect Intervals

- Update t_{start} & t_{end} for each subsequent dimension
 - If $t_1 > t_{\text{start}}$, $t_{\text{start}} = t_1$
 - If $t_2 < t_{\text{end}}$, $t_{\text{end}} = t_2$



Then Intersect Intervals


- Update t_{start} & t_{end} for each subsequent dimension
 - If $t_1 > t_{\text{start}}$, $t_{\text{start}} = t_1$
 - If $t_2 < t_{\text{end}}$, $t_{\text{end}} = t_2$



:-)


Is there an Intersection?

- If $t_{\text{start}} > t_{\text{end}}$ → box is missed




Is the Box Behind the Eyepoint?

- If $t_{\text{end}} < t_{\text{min}}$ → box is behind



Return the Correct Intersection

- If $t_{\text{start}} > t_{\text{min}}$ → closest intersection at t_{start}
- Else → closest intersection at t_{end}
 - Eye is inside box



Ray-Box Intersection Summary

- For each dimension,
 - If $R_{dx} = 0$ (ray is parallel) AND $R_{ox} < X_1$ or $R_{ox} > X_2 \rightarrow$ **no intersection**
- For each dimension, calculate intersection distances t_1 and t_2
 - $t_1 = (X_1 - R_{ox}) / R_{dx}$ $t_2 = (X_2 - R_{ox}) / R_{dx}$
 - If $t_1 > t_2$, swap
 - Maintain an interval $[t_{start}, t_{end}]$, intersect with current dimension
 - If $t_1 > t_{start}$, $t_{start} = t_1$ If $t_2 < t_{end}$, $t_{end} = t_2$
- If $t_{start} > t_{end} \rightarrow$ **box is missed**
- If $t_{end} < t_{min} \rightarrow$ **box is behind**
- If $t_{start} > t_{min} \rightarrow$ **closest intersection at t_{start}**
- Else \rightarrow **closest intersection at t_{end}**

Efficiency Issues

- $1/R_{dx}$, $1/R_{dy}$ and $1/R_{dz}$ can be pre-computed and shared for many boxes

Only difference is X_1 and X_2


Efficiency Issues

- $1/R_{dx}$, $1/R_{dy}$ and $1/R_{dz}$ can be pre-computed and shared for many boxes


Only difference is X_1 and X_2

Questions?

Bounding Box of a Triangle




Bounding Box of a Sphere



Bounding Box of a Plane


$$(x_{\max}, y_{\max}, z_{\max}) \\ = (+\infty, +\infty, +\infty)^*$$



$$(x_{\min}, y_{\min}, z_{\min}) \\ = (-\infty, -\infty, -\infty)^*$$

* unless \mathbf{n} is exactly perpendicular to an axis

Bounding Box of a Group




$$(x_{\min}, y_{\min}, z_{\min}) = (\min(x_{\min_a}, x_{\min_b}), \\ \min(y_{\min_a}, y_{\min_b}), \\ \min(z_{\min_a}, z_{\min_b}))$$

$$(x_{\max}, y_{\max}, z_{\max}) = (\max(x_{\max_a}, x_{\max_b}), \\ \max(y_{\max_a}, y_{\max_b}), \\ \max(z_{\max_a}, z_{\max_b}))$$


Bounding Box of a Transform

Bounding box of transformed object IS NOT
the transformation of the bounding box!



Bounding Box of a Transform


Bounding box of transformed object IS NOT
the transformation of the bounding box!



Questions?


Are Bounding Volumes Enough?

- If ray hits bounding volume,
must we test all primitives inside it?
 - Lots of work, think of a 1M-triangle mesh




Bounding Volume Hierarchies

- If ray hits bounding volume, must we test all primitives inside it?
 - Lots of work, think of a 1M-triangle mesh
- You guessed it already, we'll split the primitives in groups and build recursive bounding volumes




Bounding Volume Hierarchy (BVH)

- Find bounding box of objects/primitives
- Split objects/primitives into two, compute child BVs
- Recurse, build a binary tree




Bounding Volume Hierarchy (BVH)

- Find bounding box of objects/primitives
- Split objects/primitives into two, compute child BVs
- Recurse, build a binary tree




Bounding Volume Hierarchy (BVH)

- Find bounding box of objects/primitives
- Split objects/primitives into two, compute child BVs
- Recurse, build a binary tree




Bounding Volume Hierarchy (BVH)

- Find bounding box of objects/primitives
- Split objects/primitives into two, compute child BVs
- Recurse, build a binary tree




Bounding Volume Hierarchy (BVH)

- Find bounding box of objects/primitives
- Split objects/primitives into two, compute child BVs
- Recurse, build a binary tree




Where to Split Objects?

- At midpoint of current volume OR
- Sort, and put half of the objects on each side OR
- Use modeling hierarchy




Where to Split Objects?

- At midpoint of current volume OR
- Sort, and put half of the objects on each side OR
- Use modeling hierarchy




Questions?!


Ray-BVH Intersection



Ray-BVH Intersection




Ray-BVH Intersection




Early Termination

What if...




Intersection with BVH

However...




Intersection with BVH



Intersection with BVH

Early termination is powerful,
but do it cautiously



BVH Discussion

- Advantages
 - easy to construct
 - easy to traverse
 - binary tree (=simple structure)
- Disadvantages
 - may be difficult to choose a good split for a node
 - poor split may result in minimal spatial pruning


BVH Discussion

- Advantages
 - easy to construct
 - easy to traverse
 - binary tree (=simple structure)
- Disadvantages
 - may be difficult to choose a good split for a node
 - poor split may result in minimal spatial pruning
- Still one of the best methods
 - Recommended for your first hierarchy!

- Advantages
 - easy to construct
 - easy to traverse
 - binary tree (=simple structure)
- Disadvantages
 - may be difficult to choose a good split for a node
 - poor split may result in minimal spatial pruning
- Still one of the best methods
 - Recommended for your first hierarchy!

Kd-trees

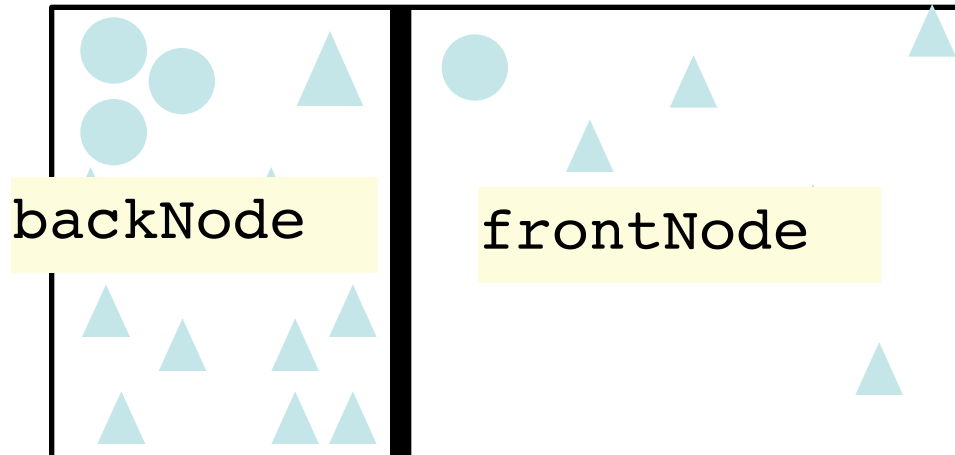
- Probably most popular acceleration structure
- Binary tree, axis-aligned splits
 - Each node splits space in half along an axis-aligned plane
- A space partition: The nodes do not overlap!
 - This is in contrast to BVHs



Data Structure

KdTreeNode:

```
KdTreeNode* backNode, frontNode //children
int dimSplit // either x, y or z
float splitDistance
    // from origin along split axis
boolean isLeaf
List of triangles //only for leaves
```




here $\text{dimSplit} = 0$ (x axis)

$X = \text{splitDistance}$


Kd-tree Construction

- Start with scene axis-aligned bounding box
- Decide which dimension to split (e.g. longest)
- Decide at which distance to split (not so easy)




Kd-tree Construction - Split

- Distribute primitives to each side
- If a primitive overlaps split plane, assign to both sides



Kd-tree Construction - Recurse

- Stop when minimum number of primitives reached
- Other stopping criteria possible



Questions?


- Further reading on efficient Kd-tree construction
 - [Hunt, Mark & Stoll, IRT 2006](#)
 - [Zhou et al., SIGGRAPH Asia 2008](#)

Zhou et al.




Kd-tree Traversal - High Level

- If leaf, intersect with list of primitives
- If intersects back child, recurse
- If intersects front child, recurse




Kd-tree Traversal, Naïve Version

- Could use bounding box test for each child
- But redundant calculation: bbox similar to that of parent node, plus axis aligned, one single split




Kd-tree Traversal, Smarter Version

- Get main bbox intersection from parent
 - t_{near} , t_{far}
- Intersect with splitting plane
 - easy because axis aligned



Kd-tree Traversal - Three Cases


- Intersects only back, only front, or both
- Can be tested by examining t , t_{start} and t_{end}



Kd-tree traversal - three cases


- If $t > t_{\text{end}} \Rightarrow$ intersect only front
- If $t < t_{\text{start}} \Rightarrow$ intersect only back

Note: “Back” and “Front” depend on ray direction!



Kd-tree Traversal Pseudocode

```
travers(orig, dir, t_start, t_end):  
    #adapted from Ingo Wald's thesis  
    #assumes that dir[self.dimSplit] > 0  
    if self.isLeaf:  
        return intersect(self.listOfTriangles, orig, dir, t_start, t_end)  
    t = (self.splitDist - orig[self.dimSplit]) / dir[self.dimSplit];  
    if t <= t_start:  
        # case one, t <= t_start <= t_end -> cull front side  
        return self.backSideNode.traverse(orig, dir, t_start, t_end)  
    elif t >= t_end:  
        # case two, t_start <= t_end <= t -> cull back side  
        return self.frontSideNode.traverse(orig, dir, t_start, t_end)  
    else:  
        # case three: traverse both sides in turn  
        t_hit = self.frontSideNode.traverse(orig, dir, t_start, t)  
        if t_hit <= t: return t_hit; # early ray termination  
        return self.backSideNode.traverse(orig, dir, t, t_end)
```



Important!

```
travers(orig, dir, t_start, t_end):
```

```
#adapted from Ingo Wald's thesis
```

```
#assumes that dir[self.dimSplit] > 0
```

```
if self.isLeaf:
```

```
    return intersect(self.listOfTriangles, orig, dir, t_start, t_end)
```

```
t = (self.splitDist - orig[self.dimSplit]) / dir[self.dimSplit];
```

```
if t <= t_start:
```

```
# case one, t <= t_start <= t_end -> cull front side
```

```
    return self.backSideNode.traverse(orig, dir, t_start, t_end)
```

```
elif t >= t_end:
```

```
# case two, t_start <= t_end <= t -> cull back side
```

```
    return self.frontSideNode.traverse(orig, dir, t_start, t_end)
```




```
else:
```

```
# case three: traverse both sides in turn
```

```
t_hit = self.frontSideNode.traverse(orig, dir, t_start, t)
```




```
if t_hit <= t: return t_hit; # early ray termination
```

```
return self.backSideNode.traverse(orig, dir, t, t_end)
```




Early termination is powerful!

```
travers(orig, dir, t_start, t_end):  
    #adapted from Ingo Wald's thesis  
    #assumes that dir[self.dimSplit] > 0  
    if self.isLeaf:  
        return intersect(self.listOfTriangles, orig, dir, t_start, t_end)  
    t = (self.splitDist - orig[self.dimSplit]) / dir[self.dimSplit];  
    if t <= t_start:  
        # case one, t <= t_start <= t_end -> cull front side  
        return self.backSideNode.traverse(orig, dir, t_start, t_end)  
    elif t >= t_end:  
        # case two, t_start <= t_end <= t -> cull back side  
        return self.frontSideNode.traverse(orig, dir, t_start, t_end)  
    else:  
        # case three: traverse both sides in turn  
        t_hit = self.frontSideNode.traverse(orig, dir, t_start, t)  
        if t_hit <= t: return t_hit; # early ray termination  
        return self.backSideNode.traverse(orig, dir, t, t_end)
```




Early termination is powerful

- If there is an intersection in the first node, don't visit the second one
- Allows ray casting to be reasonably independent of scene depth complexity



Recap: Two main gains



- Only intersect with triangles “near” the line
- Stop at the first intersection




Two main gains

```
travers(orig, dir, t_start, t_end):  
    #adapted from Ingo Wald's thesis  
    #assumes that dir[self.dimSplit] > 0  
    if self.isLeaf:  
        return intersect(self.listOfTriangles, orig, dir, t_start, t_end)  
    t = (self.splitDist - orig[self.dimSplit]) / dir[self.dimSplit];  
    if t <= t_start:  
        # case one, t <= t_start <= t_end -> cull front side  
        return self.backSideNode.traverse(orig, dir, t_start, t_end)  
    elif t >= t_end:  
        # case two, t_start <= t_end <= t -> cull back side  
        return self.frontSideNode.traverse(orig, dir, t_start, t_end)  
    else:  
        # case three: traverse both sides in turn  
        t_hit = self.frontSideNode.traverse(orig, dir, t_start, t)  
        if t_hit <= t: return t_hit; # early ray termination  
        return self.backSideNode.traverse(orig, dir, t, t_end)
```

Only near line



stop at first intersection




Important Details

- For leaves, do NOT report intersection if t is not in $[t_{\text{near}}, t_{\text{far}}]$.
 - Important for primitives that overlap multiple nodes!

Important Details

- For leaves, do NOT report intersection if t is not in $[t_{\text{near}}, t_{\text{far}}]$.
 - Important for primitives that overlap multiple nodes!




Important Details

- For leaves, do NOT report intersection if t is not in $[t_{\text{near}}, t_{\text{far}}]$.
 - Important for primitives that overlap multiple nodes!
- Need to take direction of ray into account
 - Reverse back and front if the direction has negative coordinate along the split dimension
- Degeneracies when ray direction is parallel to one axis

- For leaves, do NOT report intersection if t is not in $[t_{\text{near}}, t_{\text{far}}]$.
 - Important for primitives that overlap multiple nodes!
- Need to take direction of ray into account
 - Reverse back and front if the direction has negative coordinate along the split dimension
- Degeneracies when ray direction is parallel to one axis


Where to split for construction?

- Example for baseline
- Note how this ray traverses easily: one leaf only




Split in the Middle

- Does not conform to empty vs. dense areas
- Inefficient traversal – Not so good!



Split in the Median

- Tries to balance tree (leaf node same distance to root), but does not conform to empty vs. dense areas
- Inefficient traversal – Not good




Optimizing Splitting Planes

- Most people use the Surface Area Heuristic (SAH)
 - [MacDonald and Booth 1990, “Heuristic for ray tracing using space subdivision”, Visual Computer](#)
- Idea: simple probabilistic prediction of traversal cost based on split distance
- Then try different possible splits and keep the one with lowest cost
- Further reading on efficient Kd-tree construction
 - [Hunt, Mark & Stoll, IRT 2006](#)
 - [Zhou et al., SIGGRAPH Asia 2008](#)

Surface Area Heuristic

- Probability that we need to intersect a child
 - Area of the bbox of that child
(exact for uniformly distributed rays)
- Cost of the traversal of that child
 - number of primitives (simplistic heuristic)
- This heuristic likes to put big densities of primitives in small-area nodes



Is it Important to Optimize Splits?

- You need extra efforts in the construction
 - Sample different potential coordinates for splitting
 - Compute how many objects on the left/right.
- Given the same traversal code, the quality of Kd-tree construction can have a big impact on performance, e.g. a factor of 2 compared to naive middle split

Hard-core efficiency considerations

- See e.g. Ingo Wald's PhD thesis
 - <http://www.sci.utah.edu/~wald/PhD/index.html/>
- Calculation
 - Optimized barycentric ray-triangle intersection
- Memory
 - Make kd-tree node as small as possible (dirty bit packing, make it 8 bytes)
- Parallelism
 - Single instruction, multiple data (SIMD) extensions, trace 4 rays at a time, mask results where they disagree

Pros and Cons of Kd trees

- Pros
 - Simple code
 - Efficient traversal
 - Can conform to data
- Cons
 - costly construction, not great if you work with moving objects

Questions?

- For extensions to moving scenes, see [Real-Time KD-Tree Construction on Graphics Hardware, Zhou et al., SIGGRAPH 2008](#)



That's All For Today

- Further reading: [Shirley: Realistic Ray Tracing](#) Dutre et al.: [Advanced Global Illumination](#)